

BIRD Run Log Analysis: 20260519_061253

Inputs

- detailed_log: logs/output/20260519_061253_detailed_log.jsonl
- predictions: logs/output/20260519_061253_predictions.json
- dev: dev_20240627/dev.json
- pipeline: logs/pipeline_20260519_061253.log
- token_method: hf_tokenizer:google/gemma-4-31B-it

Overall

generation rows	evaluated questions	correct	incorrect	accuracy	finished	errors
1534	1533	1088	444	71.02%	1534	0

Timing And Tool Calls

metric	mean	median	p90	max
Generated tokens	1711.79	1578.00	2537.10	9642.00
Estimated gen seconds at 100 tok/s	17.12	15.78	25.37	96.42
Estimated gen seconds at 200 tok/s	8.56	7.89	12.69	48.21
Tool seconds	0.37	0.01	0.40	45.70
Turns per question	3.44	3.00	4.00	10.00
Attempted tool calls	2.44	2.00	3.00	9.00
Executed tool calls	1.75	2.00	2.00	8.00

By Difficulty

difficulty	generation rows	accuracy	correct	incorrect	avg tokens	avg turns	att
challenging	145	59.72%	86	58	2489.06	3.58	2.5

moderate	464	63.36%	294	170	1991.32	3.51	2.5
simple	923	76.57%	706	216	1449.44	3.38	2.3

By Database

db_id	questions	accuracy	avg turns	avg tokens	avg gen s @100	avg gen s @200
california_schools	89	73.03%	3.47	1812.64	18.13	9.06
card_games	191	65.45%	3.43	1504.75	15.05	7.52
codebase_community	186	74.05%	3.25	1356.52	13.57	6.78
debit_card_specializing	64	73.44%	3.62	2082.44	20.82	10.41
european_football_2	129	68.99%	3.42	1996.15	19.96	9.98
financial	106	72.64%	3.43	1992.78	19.93	9.96
formula_1	174	62.07%	3.48	1739.35	17.39	8.70
student_club	158	87.97%	3.60	1781.49	17.81	8.91
superhero	129	90.70%	3.36	1543.19	15.43	7.72
thrombosis_prediction	163	56.44%	3.42	1787.28	17.87	8.94
toxicology	145	64.14%	3.46	1712.47	17.12	8.56

Notes

- Attempted tool calls are counted from `<tool_call>` blocks in model output.
- Executed tool calls are turns whose action starts with `tool:` and have a recorded tool response path.
- Generated tokens are counted with the configured tokenizer when available; otherwise they use `ceil(output_chars / 4)`.
- Estimated generation seconds are per question: total generated tokens for that question divided by 100 or 200 tok/s.
- Headline accuracy comes directly from the archived predictions file.
- Per-generation-row CSV accuracy is joined from predictions by `(db_id, question)`.